

Name: Nachiket Patwardhan | student\_ID: [801208762](#)

```
In [189]: import numpy as np
import sklearn.datasets as ds
import matplotlib.pyplot as plt
from matplotlib import style
style.use('ggplot')
%matplotlib inline
import re
import pandas as pd
```

Numpy is library for scientific computing in Python. It has efficient implementation of n-dimensional array (tensor) manipulations, which is useful for machine learning applications.

```
In [190]: import numpy as np
```

We can convert a list into numpy array (tensor)

```
In [191]: b = [[1, 2, 4], [2, 6, 9]]
a = np.array(b)
a
```

```
Out[191]: array([[1, 2, 4],
                [2, 6, 9]])
```

We can check the dimensions of the array

```
In [192]: a.shape
```

```
Out[192]: (2, 3)
```

We can apply simple arithmetic operation on all element of a tensor

```
In [193]: a * 3
```

```
Out[193]: array([[ 3,  6, 12],
                [ 6, 18, 27]])
```

You can transpose a tensor

```
In [194]: print(a.T.shape)
a.T
```

```
(3, 2)
```

```
Out[194]: array([[1, 2],
                [2, 6],
                [4, 9]])
```

You can apply aggregate functions on the whole tensor

```
In [195]: np.sum(a)
```

```
Out[195]: 24
```

or on one dimension of it

```
In [196]: np.sum(a, axis=0)
```

```
Out[196]: array([ 3,  8, 13])
```

```
In [197]: np.sum(a, axis=1)
```

```
Out[197]: array([ 7, 17])
```

We can do element-wise arithmetic operation on two tensors (of the same size)

```
In [198]: c1 = np.array([[1, 2, 4], [2, 6, 9]])
          c2 = np.array([[2, 3, 5], [1, 2, 1]])
          c1 * c2
```

```
Out[198]: array([[ 2,  6, 20],
                 [ 2, 12,  9]])
```

If you want to multiply all columns of a tensor by vector (for example if you want to multiply all data features by their labels) you need a trick. This multiplication shows up in calculating the gradients.

```
In [199]: a = np.array([[1, 2, 4], [2, 6, 9]])
          b = np.array([1, -1])
          print(a)
          print(b)
```

```
[[1 2 4]
 [2 6 9]
 [ 1 -1]]
```

Here we want to multiply the first row of a by 1 and the second row of a by -1. Simply multiplying a by b does not work because a and b do not have the same dimension

```
In [200]: a * b
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-200-9bc1a869709f> in <module>
----> 1 a * b

ValueError: operands could not be broadcast together with shapes (2,3) (2,)
```

To do this multiplication we first have to assume b has one column and then repeat the column of b with the number of columns in a. We use tile function to do that

```
In [201]: b_repeat = np.tile(b, (a.shape[1],1)).T
          print(b_repeat.shape)
          b_repeat

(2, 3)
```

```
Out[201]: array([[ 1,  1,  1],
                 [-1, -1, -1]])
```

Now we can multiply each column of a by b:

```
In [202]: a * b_repeat
```

```
Out[202]: array([[ 1,  2,  4],
                 [-2, -6, -9]])
```

You can create initial random vector using numpy (using  $N(0,1)$ ):

```
In [203]: mu = 0 #mean
          sigma = 1 #standard deviation
          r = np.random.normal(mu,sigma, 1000) #draws 1000 samples from a normal distribution
```

We can apply functions on tensors


```
In [204]: #implementation of Normal distribution
          def normal(x, mu, sigma):
              return np.exp( -0.5 * ((x-mu)/sigma)**2)/np.sqrt(2.0*np.pi*sigma**2)

          #probability of samples on the Normal distribution
          probabilities = normal(r, mu, sigma)
```

Numpy has useful APIs for analysis. Here we plot the histogram of samples and also plot the probabilities to see if the samples follow the normal distribution.

```
In [205]: counts, bins = np.histogram(r,50,density=True)
plt.hist(bins[:-1], bins, weights=counts)
plt.scatter(r, probabilities, c='b', marker='.')
```

Out[205]: <matplotlib.collections.PathCollection at 0x132a0a830d0>



```
In [206]: def read_data(filename):
    f = open(filename, 'r')
    p = re.compile(',')
    xdata = []
    ydata = []
    header = f.readline().strip()
    varnames = p.split(header)
    namehash = {}
    for l in f:
        li = p.split(l.strip())
        xdata.append([float(x) for x in li[:-1]])
        ydata.append(float(li[-1]))

    return np.array(xdata), np.array(ydata)
```

Assuming our data is x is available in numpy we use numpy to implement logistic regression

```
In [207]: (xtrain_whole, ytrain_whole) = read_data('spambase-train.csv')
(xtest, ytest) = read_data('spambase-test.csv')
```

```
In [208]: print("The shape of xtrain:", xtrain_whole.shape)
print("The shape of ytrain:", ytrain_whole.shape)
print("The shape of xtest:", xtest.shape)
print("The shape of ytest:", ytest.shape)
```

```
The shape of xtrain: (3601, 54)
The shape of ytrain: (3601,)
The shape of xtest: (1000, 54)
The shape of ytest: (1000,)
```

before training make we normalize the input data (features)

```
In [209]: xmean = np.mean(xtrain_whole, axis=0)
xstd = np.std(xtrain_whole, axis=0)
xtrain_normal_whole = (xtrain_whole-xmean) / xstd
xtest_normal = (xtest-xmean) / xstd
```

We need to create a validation set. We create an array of indices and permute it.

```
In [210]: permute_indices = np.random.permutation(np.arange(xtrain_whole.shape[0]))
```

We keep the first 2600 data points as the training data and rest as the validation data

```
In [211]: xtrain_normal = xtrain_normal_whole[permute_indices[:2600]]
ytrain = ytrain_whole[permute_indices[:2600]]
xval_normal = xtrain_normal_whole[permute_indices[2600:]]
yval = ytrain_whole[permute_indices[2600:]]
```

Initializing the weights and bias with random values from  $N(0,1)$

```
In [212]: weights = np.random.normal(0, 1, xtrain_normal.shape[1]);
bias = np.random.normal(0,1,1)
```

```
In [213]: #the sigmoid function
def sigmoid(z):
    s = 1 / (1 + np.exp(-z))
    return s
```

We can use dot-product from numpy to calculate the margin and pass it to the sigmoid function

```
In [214]: #w: weight vector (numpy array of size n)
#b: numpy array of size 1
#returns p(y=1/x, w, b)
def prob(x, w, b):
    z = sigmoid(np.dot(x,w) + b)
    #print(z)
    return z
```

You can also calculate  $l_2$  penalty using linalg library of numpy

```
In [215]: np.linalg.norm(weights)
```

```
Out[215]: 7.412178281536248
```

$$\text{Cross Entropy Loss} = - \sum_{(y^i, \mathbf{x}^i) \in \mathcal{D}} y^i \log p(y = 1 | \mathbf{x}^i; \mathbf{w}, b) + (1 - y^i) \log(1 - p(y = 1 | \mathbf{x}^i; \mathbf{w}, b)) -$$

```
In [216]: #w: weight vector (numpy array of size n)
#y_prob: p(y|x, w, b)
#y_true: class variable data
#lambda_: L2 penalty coefficient
#returns the cross entropy loss
def loss(w, y_prob, y_true, lambda_):
    Cross_Entropy_loss = -(1) * np.sum((y_true * np.log10(y_prob)) + ((1 - y_true
    return Cross_Entropy_loss
```

```
In [217]: #x: input variables (data of size m x n with m data point and n features)
#w: weight vector (numpy array of size n)
#y_prob: p(y|x, w, b)
#y_true: class variable data
#lambda_: L2 penalty coefficient
#returns tuple of gradient w.r.t w and w.r.t to bias

def grad_w_b(x, w, y_prob, y_true, lambda_):
    grad_w = np.zeros(w.shape)
    #m = len(x)
    grad_b = 0.0
    grad_w = np.dot(x.T, (y_prob - y_true)) + (2 * lambda_ * w)
    grad_b = np.sum(y_prob - y_true)

    return(grad_w, grad_b)
```

In [224]:

```

#lambda_ is the coeffienct of l2 norm penalty
#learning_rate is Learning rate of gradient descent algorithm
#max_iter determines the maximum number of iterations if the gradients descent do
#continue the training while gradient > 0.1 or the number steps is Less max_iter

#returns model as tuple of (weights,bias)

def fit(x, y_true, learning_rate, lambda_, max_iter, verbose=0):
    weights = np.random.normal(0, 1, x.shape[1]);
    bias = np.random.normal(0,1,1)
    #raise NotImplementedError
    costs = []
    i = 0
    while(i < max_iter):
        #To calculate y_prob
        y_prob = prob(x, weights, bias)

        Cross_Entropy_loss = loss(weights, y_prob, y_true, lambda_)

        #costs = np.append(cost)

        grad_w, grad_b = grad_w_b(x, weights, y_prob, y_true, lambda_)

        #print(grad_w)
        #print(grad_b)

        #Condition to break if grad_w and grad_b is less than 0.1
        if((np.linalg.norm(grad_w) < 0.1) and (np.linalg.norm(grad_b) < 0.1)):
            break
        else:
            weights = weights - learning_rate * grad_w

            bias = bias - learning_rate * grad_b

        i = i+1

    return (weights, bias)

```

In [225]:

```

def accuracy(x, y_true, model):
    w, b = model
    return np.sum((prob(x, w, b)>0.5).astype(float) == y_true) / y_true.shape[0]

```

In [226]:

```

learning_rate = 0.001
lambda_ = 1.0

model = fit(xtrain_normal, ytrain, learning_rate, lambda_, 10000, verbose=1)

```

In [227]:

```

print("Train accuracy: ", accuracy(xtrain_normal, ytrain, model))

```

Train accuracy: 0.9338461538461539

In [228]: *#grid search for finding the best hyperparams and model*

```
best_model = None
best_val = -1
for lr in [0.01, 0.001, 0.0001, 0.00001]:
    for la in [5, 2, 1, 0.1, 0.01]:
        model = fit(xtrain_normal, ytrain, lr, la, 10000)
        val_acc = accuracy(xval_normal, yval, model)
        print(lr, la, val_acc)
        if val_acc > best_val:
            best_val = val_acc
            best_model = model
```

<ipython-input-216-e5e81e887299>:7: RuntimeWarning: divide by zero encountered in log10

```
Cross_Entropy_loss = -(1) * np.sum((y_true * np.log10(y_prob)) + ((1 - y_true) * np.log10(1 - y_prob))) + ((lambda_/2) * (w**2))
```

<ipython-input-216-e5e81e887299>:7: RuntimeWarning: invalid value encountered in multiply

```
Cross_Entropy_loss = -(1) * np.sum((y_true * np.log10(y_prob)) + ((1 - y_true) * np.log10(1 - y_prob))) + ((lambda_/2) * (w**2))
```

```
0.01 5 0.7922077922077922
0.01 2 0.903096903096903
0.01 1 0.9150849150849151
0.01 0.1 0.9050949050949051
0.01 0.01 0.8561438561438561
0.001 5 0.9300699300699301
0.001 2 0.9330669330669331
0.001 1 0.9340659340659341
0.001 0.1 0.9340659340659341
0.001 0.01 0.9340659340659341
0.0001 5 0.9300699300699301
0.0001 2 0.9330669330669331
0.0001 1 0.9340659340659341
0.0001 0.1 0.9340659340659341
0.0001 0.01 0.9340659340659341
1e-05 5 0.9300699300699301
1e-05 2 0.9330669330669331
1e-05 1 0.9340659340659341
1e-05 0.1 0.936063936063936
1e-05 0.01 0.9340659340659341
```

In [229]: `print("Test accuracy: ", accuracy(xtest_normal, ytest, best_model))`

Test accuracy: 0.941

In [ ]:



